



- Rcvbuffer → typically 4096 bytes
- remaining buffer space updated in receiver's window (rwnd) field
- sender limits unACKed data to rwnd to ensure no buffer overflow

$$\text{LastByteRead} - \text{LastByteRead} \leq \text{Rcvbuffer}$$

$$\text{rwnd} = \text{Rcvbuffer} - [\text{LastByteRead} - \text{LastByteRead}]$$

event	response
→ pkt w/ expected seq #: all data upto this already ACKed	delayed ACK: wait 500ms for next seg
→ pkt w/ expected seq # one other pending	immediate cumulative ACK
→ out of order higher than expect. gap detected	immediate send duplicate ACK
→ seg that partially/fully fills gap	immediate send ACK for this, assuming it fills gap from lower end

Note: Fast retransmit
3 additional duplicate ACKs for a seg
↓
retransmit without waiting for timeout

sender creates segment w/ seq # for application data; start timer if not already running for oldest unACKed segment

timeout: retransmit, restart timer

receive ACK: if ACK is for unACKed seg:
• update status
• start timer for next unACKed seg

Transmission Control Protocol

- point to point: single sender, receives
- full duplex communication
- flow control
- connection oriented
- cumulative ACKs
- pipelining: congestion and flow control determine window size
- inorder byte stream: no message boundaries

sender port #	receiver port #
sequence number	
acknowledgement no.	
length of TCP header	
seq num C E V A P R S F receive window	
options (variable length)	
application data (variable length)	

byte stream no. of first byte in this segment's data
seq # of next byte expected from other side. seq(A) = this is ACK for flow control: no of bytes received willing to accept

TCP Congestion Control

AIMD

Additive Increase, Multiplicative Decrease

- ① Receive ACK ⇒ pkt received
↓
increase speed
- ② Timeout ⇒ pkt lost
↓
decrease speed

Slow Start

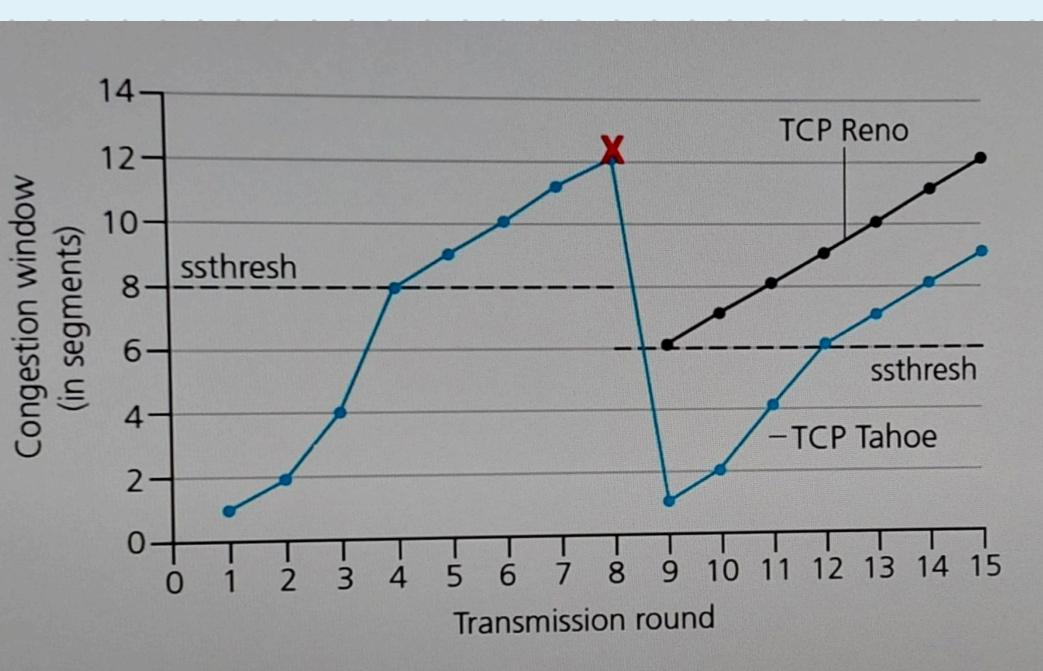
Initially cwnd = 1 MSS
Double cwnd every RTT

Done by incrementing for every ACK

Congestion Avoidance

Switch exponential increase to linear increase when cwnd gets to $\frac{1}{2}$ its value before timeout.

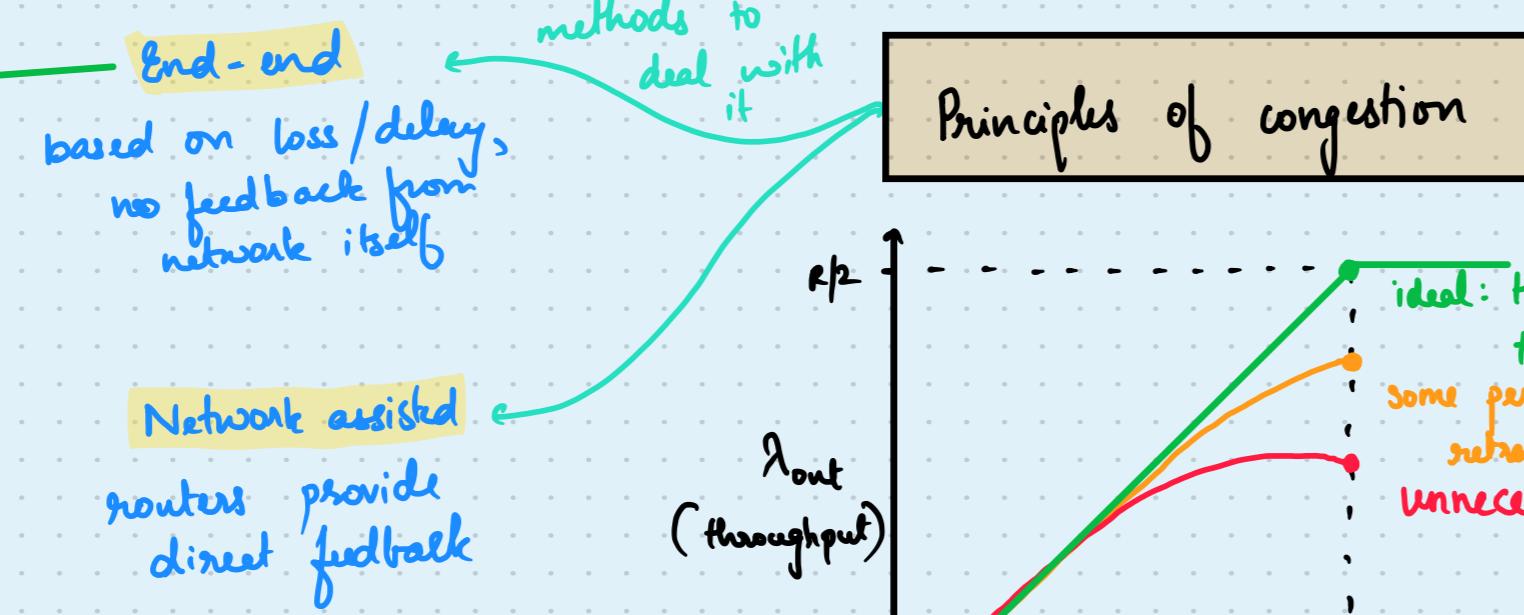
Implemented by ssthresh: set to half of cwnd on a loss event



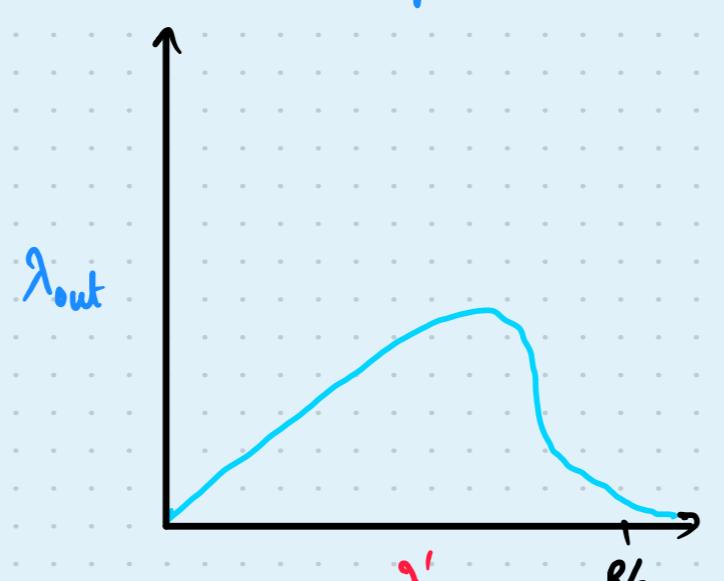
Multiplicative Decrease

3 duplicate ACKs → cut cwnd to half → TCP Reno

Timeout → set cwnd to only 1 MSS → TCP Tahoe



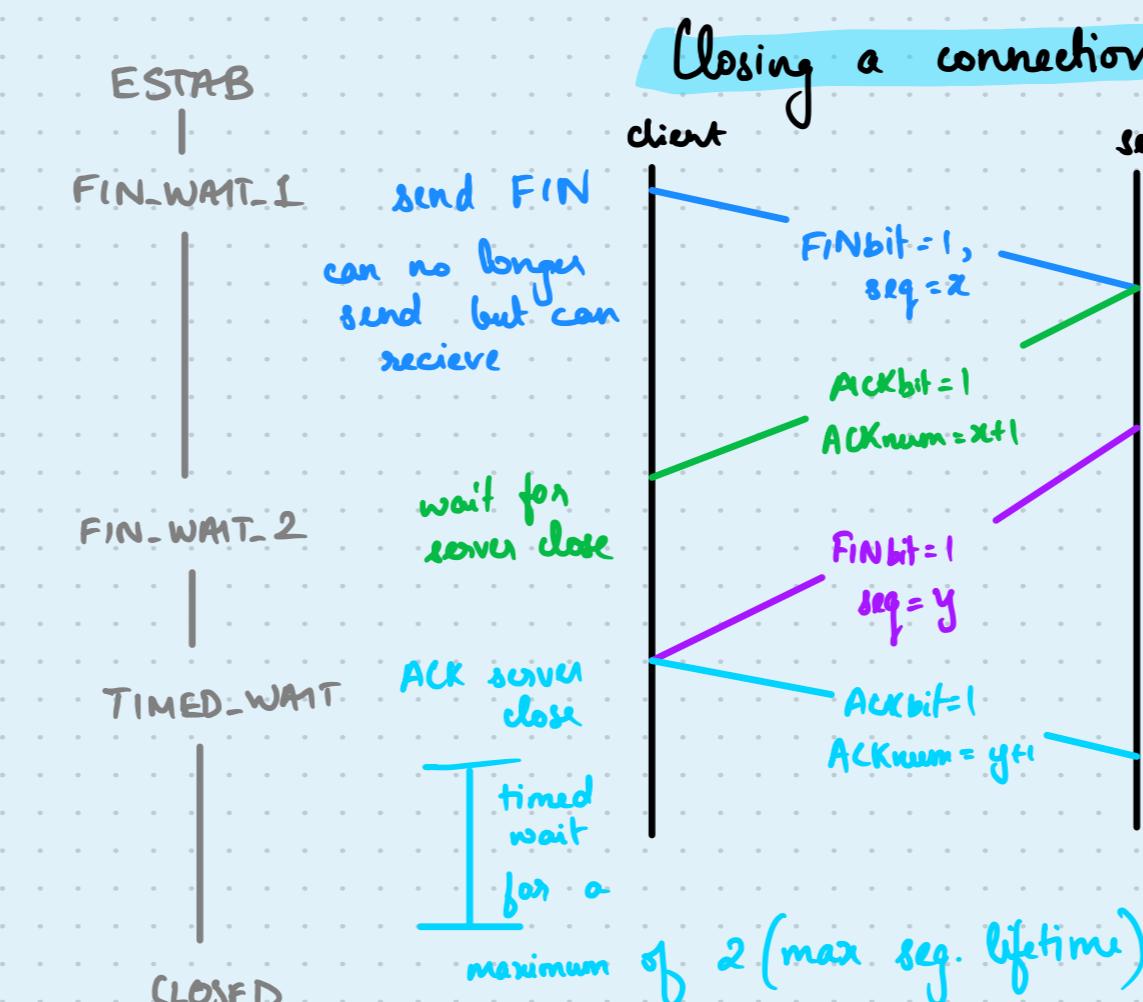
For multiple senders (red/blue):



$\lambda'_in \uparrow \Rightarrow \lambda_{out} \downarrow$

$\lambda'_in \rightarrow \lambda/2 \Rightarrow \lambda_{out} \rightarrow 0$

Any upstream transmission/buffering → wasted!



ESTAB

FIN_WAIT_1

FIN_WAIT_2

TIMED_WAIT

CLOSED

client

server



LISTEN

SYN

ESTAB

client

server



SYN-ACK

ACK

ESTABLISHED

timeout > RTT, but how to set?

exponential weighted moving avg (EWMA)

$$\text{Estimated RTT} = \alpha (\text{Estimated RTT}_{\text{old}}) + (1-\alpha)(\text{Sample RTT})$$

where typical $\alpha = 0.125$

$$\text{Timeout Interval} = \text{Estimated RTT} + 4(\text{DevRTT})$$

where DevRTT → Deviation RTT is given by EWMA of sample deviation from estimate

$$\text{DevRTT} = \beta (\text{DevRTT}_{\text{old}}) + (1-\beta)|\text{EstimatedRTT} - \text{SampleRTT}|$$

where typical $\beta = 0.25$