

Overall Process

- Scanning: Identify query keywords, tokens
- Parsing: Verify syntax
- Validation: Verify all referenced relations etc. exist
- Query tree or DAG created
- Query optimiser: Execution plan devised
- Code generator: Low level code for execution plan generated
- Runtime database processor: Low level code executed

Decomposition into Blocks

- Query decomposed into basic SELECT-FROM-WHERE blocks
- May have HAVING/GROUP BY
- These units -> translated into relational algebra equivalent

Example :

```
SELECT Lname, Fname  
FROM EMPLOYEE  
WHERE Salary > ( SELECT MAX (Salary)  
                  FROM   EMPLOYEE  
                  WHERE Dno=5 );
```

Inner Block :

```
( SELECT MAX (Salary)  
      FROM   EMPLOYEE  
      WHERE Dno=5 )
```

Outer Block :

```
SELECT Lname, Fname  
FROM EMPLOYEE  
WHERE Salary > c
```

Inner Block translated into :

$$\exists_{\text{MAX Salary}} (\sigma_{Dno=5}(\text{EMPLOYEE}))$$

Outer Block translated into :

$$\pi_{\text{Lname}, \text{Fname}} (\sigma_{\text{Salary} > c}(\text{EMPLOYEE}))$$

Query optimiser chooses execution plan **for each block**.

Semi-join, Anti-join

- Exists to be more computationally efficient than doing an entire inner join

Semi

T1.x S= T2.y

- Returns the first row of T1 that matches T1.x and T2.y

Anti

T1.x A= T2.y

- Returns the first row of T1 that does not match T1.x and T2.y

Implementing SELECT

Linear Search - Brute Force

Load every disk block one by one into main mem -> search through every record in that disk block and check condition

Binary Search

- Cond: Equality on key attribute using which the file is ordered
- More efficient than linear search

Primary Index

- Cond: Equality on key attribute that is used as primary index for file contents
- Find index -> get pointer to relevant row
- Retrieves single record at most

Hash Key

- Cond: Equality on key attribute with hash key

- $H(key_val) \rightarrow$ direct index
- Retrieves single record at most

Primary Index for Multiple Records

- Cond: $<$, $>$, \leq , \geq on key attribute that is used as primary index for file contents
- Find index \rightarrow relevant row
- Select all subsequent/preceding rows as well depending on operation

Implementing JOIN

Two-way / multiway joins

Nested Loop Join (Brute Force)

- Manually check against every possible value in the inner loop
- ```
Student.Did = Dept.Did
For every record in Student:
 Go through every record in Department:
 if Student.Did == Department.Did:
 return joined record
```

### Index-Based Nested Loop Join

- Use an index to look up value in inner loop instead of searching
- Suppose index exists on Department.Did
 

```
For every record in Student:
 matched_dept = index_lookup(Student.record.DiD)
 for each record in matched_dept:
 return joined record
```

## Sort-Merge Join

- Join on attributes on which tables are ordered
- Merge = scan both tables  $\rightarrow$  compare ID (or whatever attr) sequentially  $\rightarrow$  merge records with matching IDs and output

## Partition-Hash Join

- Two phases: partitioning and probing
- Hashing done using same hashing function on both tables
- Partitioning:
  - Single pass on the smaller table
  - Collection of records with the same value of  $H(A)$   $\rightarrow$  same hash bucket/partition
  - We assume that after partitioning, every individual bucket can fit into main memory
- Probing:
  - Pass on other table and hash value
  - Probe the hash bucket corresponding to this hash value

- Combine record with matching records from the partition

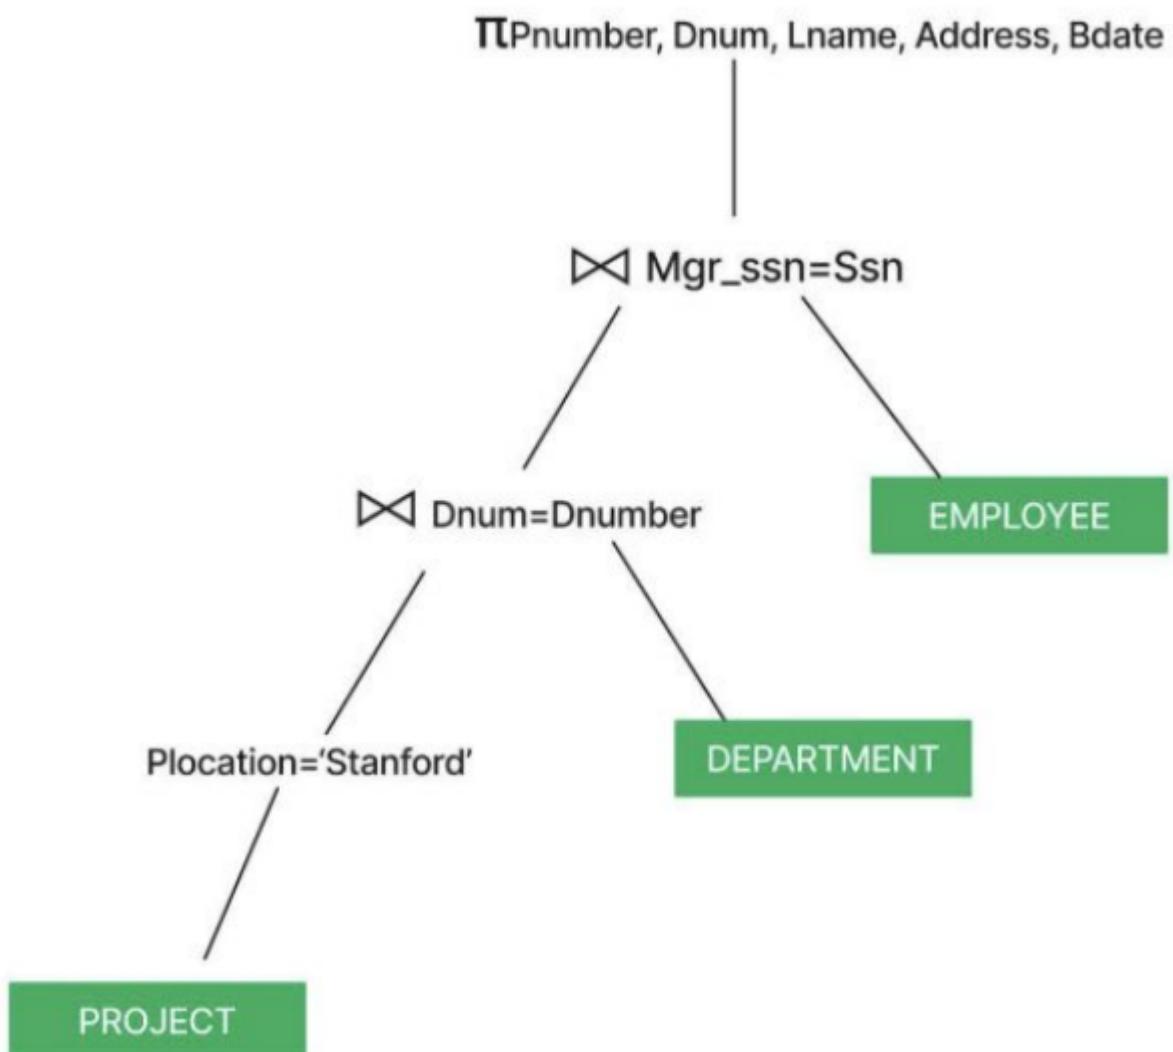
# Query Optimisation

High level SQL query  $\rightarrow$  query tree/graph  $\rightarrow$  heuristics to remove unnecessary operations  $\rightarrow$  calculate cost of each operation  $\rightarrow$  cost model compares plans  $\rightarrow$  finally selected plan is executed

# Query Tree

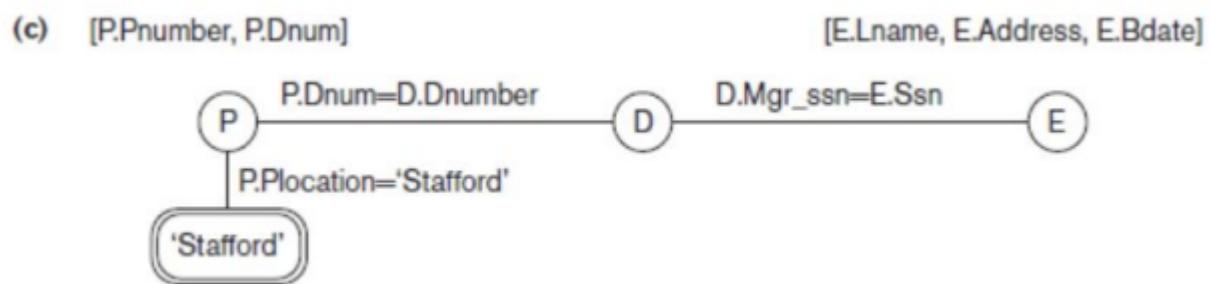
- Also called Query Evaluation/Execution Tree
  - Leaf nodes = relations
  - Internal nodes = relational algebra operators
  - Root node = output

```
SELECT P.Pnumber, P.Dnum, E.Lname, E.Address, E.Bdate
FROM PROJECT P, DEPARTMENT D, EMPLOYEE E
WHERE P.Dnum=D.Dnumber AND D.Mgr_ssn=E.Ssn AND
P.Plocation= 'Stafford';
```



# Query Graphs

- Query tree preferred since it shows a definite order
- Relations = Circles
- Constants = Double circles
- Select/Join conditions = edges
- Attributes to be retrieved = written in square brackets



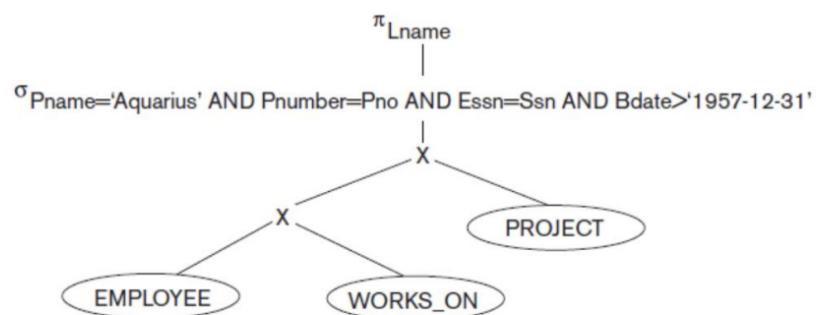
## Query Transformation

- Use heuristics to optimise query execution

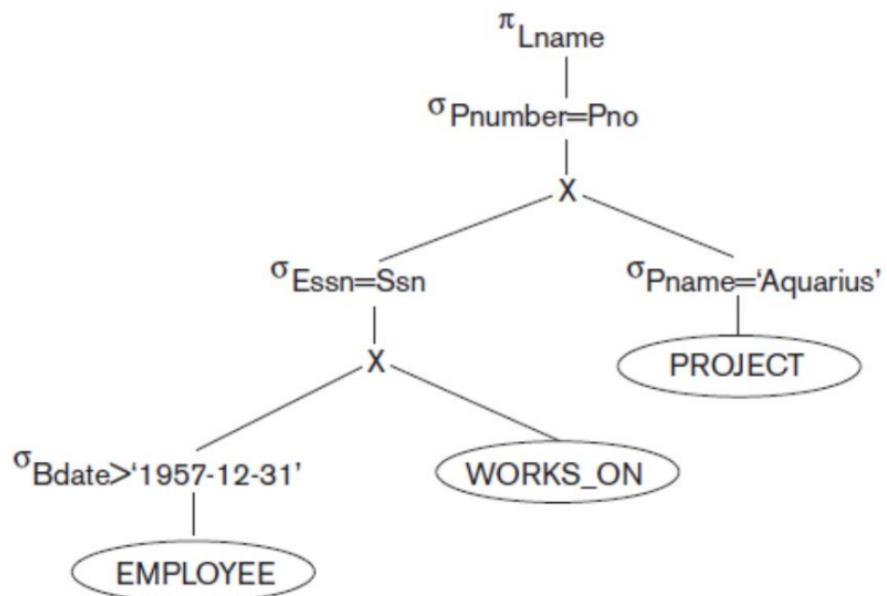
Consider the query Q

Q: `SELECT E.Lname  
FROM EMPLOYEE E, WORKS_ON W, PROJECT P  
WHERE P.Pname='Aquarius' AND P.Pnumber=W.Pno AND E.Essn=W.Ssn  
AND E.Bdate > '1957-12-31';`

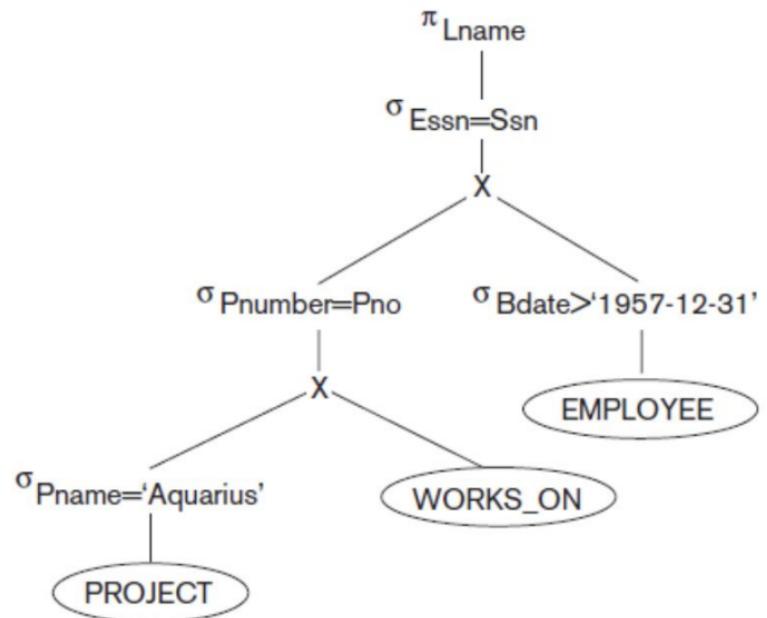
Step 1. Initial (canonical) query tree for SQL query Q.



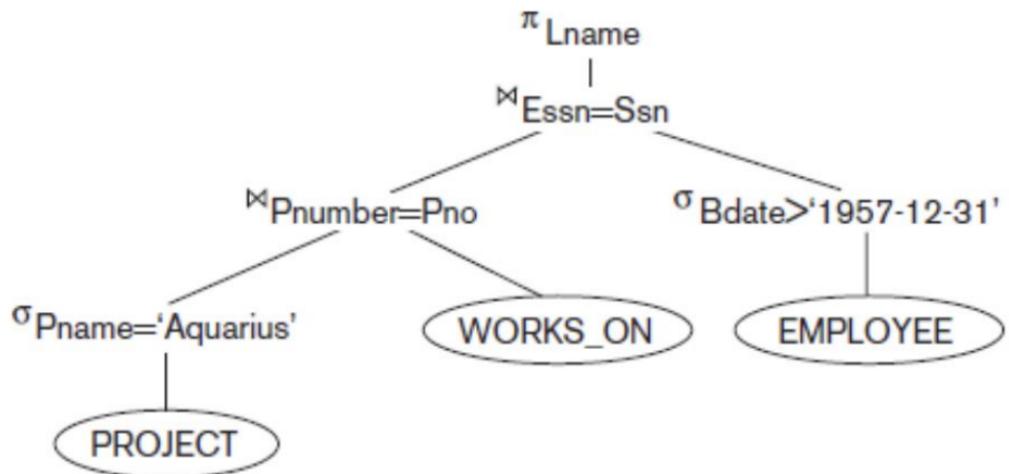
## Step 2. Moving SELECT operations down the query tree.



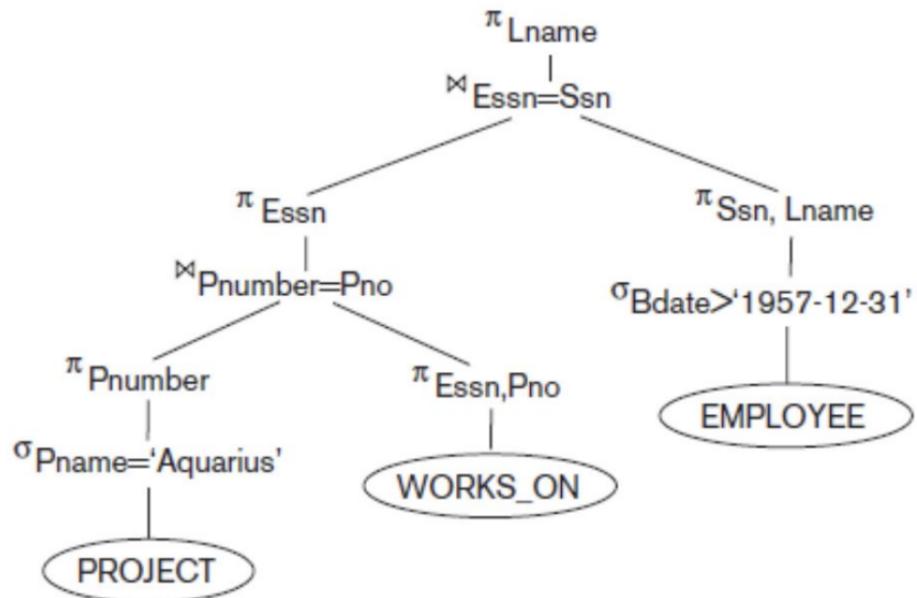
## Step 3. Applying the more restrictive SELECT operation first.



## Step 4. Replacing CARTESIAN PRODUCT and SELECT with JOIN operations.



## Step 5. Moving PROJECT operations down the query tree.



## Types of Query Optimisation

### 1. Heuristics Based

- Defined rules used to refine the query tree structure
  - Execute SELECT and PROJECT early
  - Replace Cartesian product + SELECT with JOIN
  - Execute more restrictive filters first, etc.
- No cost estimation, only heuristics

### 2. Cost-Based

- Calculate a cost for each query execution plan (estimate)
- Pick the one with the lowest one
- Use statistics from the DBMS catalog: total relations, file size, histogram etc

### 3. Semantic Query Optimisation

- Rewrite query to be more efficient
- Utilise semantic information derived from schema and integrity constraints
- Uses schema rules/constraints to eliminate redundant operations

### 4. Dynamic Query Optimisation

- Changes plan at execution
- Adjusts join order, access paths, memory usage dynamically
- Uses runtime feedback loops to refine plans

## Indexing

```
CREATE INDEX idx_salary ON Employee(salary);
EXPLAIN SELECT * FROM Employee WHERE salary > 50000;
```

```
match(p:Movie) where p.name="Movie3" detach delete p
```